

# RELIA: Accelerating the Analysis of Cloud Access Control Policies

Dan Wang<sup>1</sup>, Peng Zhang<sup>1</sup>, Zhenrong Gu<sup>1</sup>, Weibo Lin<sup>2</sup>, Shibiao Jiang<sup>2</sup>,  
Zhu He<sup>2</sup>, Xu Du<sup>2</sup>, Longfei Chen<sup>2</sup>, Jun Li<sup>2</sup>, and Xiaohong Guan<sup>1</sup>



## OVERVIEW

- Cloud providers offer flexible access control by **access control policies** to secure user's cloud resources. However, configuring access control policies is **error-prone**. Cloud providers have developed SMT-based tools to formally analyze the user-defined policies.
- Unfortunately, these analyzers are **slow**, due to the **complex regular expression** matching conditions in policies.
- We propose **RELIA**, a general method to speed up the analysis of cloud access control policies. RELIA **pre-computes** a set of **String Equivalence Classes (SECs)** based on the regular expressions in a policy, assign a **unique integer** to each SEC, and **rewrite** the regular constraints into equivalent **integer** constraints, which are easier to solve. We implement RELIA as a **transparent layer** between policy analyzer and off-the-shelf SMT solvers.
- Based on **real policies** from Huawei, we show that: when enabling RELIA, our in-house portfolio solver can **speed up** the analysis process for nearly 95% of all cases, with an average speedup of **8.21x**.

## BACKGROUND

As **cloud services thrive, resources on cloud need to be secured**. Many companies outsource their services to the cloud. To secure those assets, cloud service providers (CSPs) enable users to configure access control policies. However, these policies are complex and error-prone. A misconfiguration in the policies can lead to security risks, such as exposing private resources to the public, or allowing attackers to gain unauthorized access by escalating privileges.


**SMT-Based access analyzers are slow.** Evaluation results show that general-purpose SMT solvers (e.g., Z3, CVC4, and CVC5) are slow when analyzing some complex policies.

**String search space is huge, slowing down SMT solving.** However, a lot of strings satisfy the same set of regular constraints and thus can be merged into an equivalence class (*or string equivalence class, SEC*). We found that the total number of SECs is small for cloud access control policies: on 506 real policies from Huawei Cloud, the total number of SECs is always less than 12.


We propose **RELIA** to utilize the reduced number of SECs and accelerate string solving in policy analysis.

## KEY IDEAS

RELIA converts constraints over strings to constraints over **String Equivalence Classes**!




- Eliminate concatenations.** ACP often contains compound string variables, e.g. urn (*uniform resource name*) = org (*organization*) + ":" + id. If there is a regular expression constraint on it, it also constrains the component variables, i.e. org and id. RELIA distributes the constraints to its component constraints to achieve equisatisfiability.
- Convert regular constraints.** RELIA converts regular constraints of string variables into linear integer arithmetic (LIA) constraints over integer variables, base on which SECs satisfies the original constraint.



## HANDLING STRING CONCATENATIONS

Compound string variables in ACP have easy forms. For clarity, compound string variables in ACP have:

- Anchor points.** The component variables of a compound string variable are separated by string literals like ":" or "/", and they should not appear in these variables to eliminate ambiguity.
- One-and-only-one concatenation.** A compound variable can have multiple concatenation schemes (such as urn can be org + "user:" + user\_id for individuals, or org + "agency:" + session\_id for agencies); but one and only one concatenation scheme is valid given a value of the variable.




We can split regular constraints easily on anchor points of their automata. If compound variable s have regular constraint  $s \in R$ , it has component variables  $s_0, s_1, \dots, s_m$ , and having N split cases, It can be converted into:


$$s \in R \implies \bigvee_{i \in [1, N]} \left( \bigwedge_{j \in [1, m]} s_j \in R_i^j \right)$$

## CONVERTING REGULAR CONSTRAINTS

**Computes SECs.** RELIA computes the SECs regarding to all the regular constraints of each variable. If a variable has 2 regular constraints: fall125/. $*$  and fall125/stu/. $*$ , they divide the string space into 3 SECs, labeled:



SEC is computed by iterate through regular constraints and compute FA / set intersection and difference with previous SECs.




Finally turn regular constraints into LIA constraints. If  $s \in R$  iff  $s$  is in SECs labeled  $I_R = \{i_1, i_2, \dots, i_n\}$ , then  $s$  can be represented with integer variable  $i_s$ , and  $s \in R$  is converted into:

$$s \in R \implies \bigvee_{n \in I_R} i_s = n$$


## EVALUATIONS

RELIA is able to:


- Accelerate real policy analysis by **8.21x**;
- Analyze **8/10 harder policies** in 3s which Z3, Cvc4 and Cvc5 cannot solve most in hours;
- Generalize to and accelerate **other CSP's ACP analysis**;
- Solve **general string SMT problems** at great speed.




Speedup of real Huawei policy analysis



Speedup of difference analyses types



Speedup of translated AWS policies from Quacky



Speedup of general string SMT problems from AutomatArk



Presentation Paper